

Introduction to Data Science

HW3

Report

This report is a combination of HW2 analysis and HW3 analysis, with HW3 focusing on data cleaning. This report describes the analysis of the dataset from the CSV file "2023 June Unemployment Rate by County (Percent).csv." To guarantee data quality and completeness, the dataset was loaded from a CSV file, and different data analysis and manipulation operations were performed.

Code Summary:

The provided code accomplishes the following tasks:

Data Loading: The code starts by importing necessary Python libraries such as NumPy, Matplotlib, Pandas, and Seaborn. These libraries are frequently employed in data analysis and visualization.

The dataset is loaded into a Pandas DataFrame named 'dataset' from the file '2023 June Unemployment Rate by County (Percent).csv'. This first step is critical for gaining access to and modifying the data.

▼ Importing the libraries

```
✓ [48] import numpy as np
0s      import matplotlib.pyplot as plt
      import pandas as pd
      import seaborn as sns
```

▼ Importing the dataset

```
✓ [49] file_path = '2023 June Unemployment Rate by County (Percent).csv'
0s      dataset = pd.read_csv(file_path)
      x = dataset.iloc[:, :-1].values
      y = dataset.iloc[:, -1].values
```

Data Splitting:

It divides the dataset into two arrays, x, and y, where x contains all except the last column and y contains the last column. This is commonly done to differentiate between features (independent variables) and the desired variable (dependent variable).

```
[50] print(x)

[['Series ID' 'Region Name' 'Region Code']
 ['ALAUTA1URN' 'Autauga County, AL' '1001']
 ['ALBALD0URN' 'Baldwin County, AL' '1003']
 ...
 ['WYUINT1URN' 'Uinta County, WY' '56041']
 ['WYWASH3URN' 'Washakie County, WY' '56043']
 ['WYWEST5URN' 'Weston County, WY' '56045']]

[51] print(y)

['01-06-2023' '2.3' '2.3' ... '3.4' '3.4' '2.2']
```

Data Exploration:

It shows the first ten rows of the Data Frame, providing an overview of the data's structure and content.

```
✓ [61] print("First few rows of the DataFrame:")
0s    print(dataset.head(10))
```

First few rows of the DataFrame:

	2023 June Unemployment Rate by County (Percent)	Unnamed: 1 \
0	Series ID	Region Name
1	ALAUTA1URN	Autauga County, AL
2	ALBALD0URN	Baldwin County, AL
3	ALBARB5URN	Barbour County, AL
4	ALBIBB7URN	Bibb County, AL
5	ALBLOU9URN	Blount County, AL
6	ALBULL1URN	Bullock County, AL
7	ALBUTL3URN	Butler County, AL
8	ALCALH5URN	Calhoun County, AL
9	ALCHAM7URN	Chambers County, AL

	Unnamed: 2	Unnamed: 3
0	Region Code	01-06-2023
1	1001	2.3
2	1003	2.3
3	1005	5
4	1007	2.9
5	1009	2.3
6	1011	2.7
7	1013	3.2
8	1015	3
9	1017	2.6

Displays the last 5 rows using `dataset.tail(5)` to check the data's end.

```

0 s 0 s ## View the last few rows of the data frame
print(dataset.tail(5))

2023 June Unemployment Rate by County (Percent)      Unnamed: 1 \
3140 WYSWEE7URN Sweetwater County, WY
3141 WYTETO9URN Teton County, WY
3142 WYUINT1URN Uinta County, WY
3143 WYWASH3URN Washakie County, WY
3144 WYWEST5URN Weston County, WY

      Unnamed: 2 Unnamed: 3
3140      56037      3.6
3141      56039      1.7
3142      56041      3.4
3143      56043      3.4
3144      56045      2.2

```

Retrieves the column names using `dataset.columns` to understand the variables included in the dataset.

Generates summary statistics using `dataset.describe()`, which provides insights into the central tendencies and distributions of numerical columns.

Uses `dataset.info()` to obtain information about data types and non-null counts.

```

0 s [7] # Use the type() function to determine the data type of each object
print("Type of the data:", type(dataset))

Type of the data: <class 'pandas.core.frame.DataFrame'>

0 s [8] #Getting the dimension of the data
print(dataset.shape)

(3145, 4)

0 s 0 s #Number of Rows and Columns:
num_rows, num_columns = dataset.shape
print(f"Number of rows: {num_rows}")
print(f"Number of columns: {num_columns}")

Number of rows: 3145
Number of columns: 4

0 s [10] #Column Names:
column_names = dataset.columns
print("Column names:", column_names)

Column names: Index(['2023 June Unemployment Rate by County (Percent)', 'Unnamed: 1',
                    'Unnamed: 2', 'Unnamed: 3'],
                  dtype='object')

```

```

0 s [60] print("\nSummary statistics:")
print(dataset.describe())

Summary statistics:
2023 June Unemployment Rate by County (Percent)      Unnamed: 1 \
count      3145      3145
unique      3145      3142
top      Series ID Hancock County, KY
freq      1      2

      Unnamed: 2 Unnamed: 3
count      3145      3140
unique      3142      91
top      21091      3.1
freq      2      141

```

```

✓ [12] #Check the structure of the data frame
os print(dataset.info())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3145 entries, 0 to 3144
Data columns (total 4 columns):
 #   Column                                                                 Non-Null Count  Dtype  
---  -
 0   2023 June Unemployment Rate by County (Percent)                 3145 non-null  object 
 1   Unnamed: 1                                                         3145 non-null  object 
 2   Unnamed: 2                                                         3145 non-null  object 
 3   Unnamed: 3                                                         3140 non-null  object 
dtypes: object(4)
memory usage: 98.4+ KB
None

✓ [13] # Use the For loop to get unique value in all the columns
os for col_name in dataset.columns:
    #print(data[col_name].unique())
    print("Unique values in column", col_name, ":", dataset[col_name].unique() , "\n")

Unique values in column 2023 June Unemployment Rate by County (Percent) : ['Series ID' 'ALAUTA1URN' 'ALBALD01' 'WYWEST5URN']

Unique values in column Unnamed: 1 : ['Region Name' 'Autauga County, AL' 'Baldwin County, AL' ... 'Uinta County, WY' 'Washakie County, WY' 'Weston County, WY']

Unique values in column Unnamed: 2 : ['Region Code' '1001' '1003' ... '56041' '56043' '56045']

Unique values in column Unnamed: 3 : ['01-06-2023' '2.3' '5' '2.9' '2.7' '3.2' '3' '2.6' '2.5' '4.3' '5.2'
'2.4' '3.1' '3.9' '2.8' '2.2' '5.7' '3.7' '4.4' '3.8' '3.6' '2.1' '6.2'
'3.5' '2' '4.5' '7.5' '3.3' '10.1' '7.2' '5.1' '4' '4.9' '8.1' '10' '6.4'
'5.8' nan '18.1' '4.7' '6.3' '8.9' '8.8' '4.8' '5.5' '4.6' '7.3' '4.2'
'14.3' '5.3' '4.1' '3.4' '7.1' '5.4' '12.2' '5.6' '6.5' '16.9' '8.4'
'7.7' '9.5' '6' '5.9' '6.8' '7.6' '6.6' '1.7' '1.9' '1.8' '6.1' '8.5'
'8.2' '9.4' '6.9' '7.4' '7.8' '9.7' '9.1' '6.7' '10.2' '1.5' '1.6' '1.4'
'7.9' '14.1' '9.8' '1.2' '9' '8' '0.9' '0.4' '9.6']

```

Missing Values Handling:

Missing values are common in real-world datasets and handling them is critical to ensuring that subsequent analyses are accurate. To handle missing values, the code does the following:

Uses 'dataset.isnull().sum()' to check for missing values and reports the number of missing values in each column.

To highlight the pattern of missing values across columns, a heatmap built with Seaborn's `sns.heatmap()` is used to visualize the missing data.

```

[14] # Count rows with complete cases (no missing values)
complete_cases_count = dataset.dropna().shape[0]
# Count rows with missing values
missing_cases_count = dataset.shape[0] - complete_cases_count
# Print the counts
print("Rows with complete cases:", complete_cases_count)
print("Rows with missing values:", missing_cases_count)

Rows with complete cases: 3140
Rows with missing values: 5

[15] # Check for missing values and print the count of missing values
missing_values = dataset.isnull().sum()
print("\nMissing values:")
print(missing_values)

Missing values:
2023 June Unemployment Rate by County (Percent)    0
Unnamed: 1                                           0
Unnamed: 2                                           0
Unnamed: 3                                           5
dtype: int64

[16] missing_value = ["N/a", "na", np.nan]
dataset = pd.read_csv("2023 June Unemployment Rate by County (Percent).csv", na_values = missing_value)

[17] dataset.isnull().sum()

2023 June Unemployment Rate by County (Percent)    0
Unnamed: 1                                           0
Unnamed: 2                                           0
Unnamed: 3                                           5
dtype: int64

[18] dataset.isnull().any()

2023 June Unemployment Rate by County (Percent)    False
Unnamed: 1                                           False
Unnamed: 2                                           False
Unnamed: 3                                           True
dtype: bool

```

```

missing_positions = missing_positions = np.where(pd.isna(dataset["Unnamed: 3"]))
print("Positions of missing values in variable Unnamed: 3 :",missing_positions)

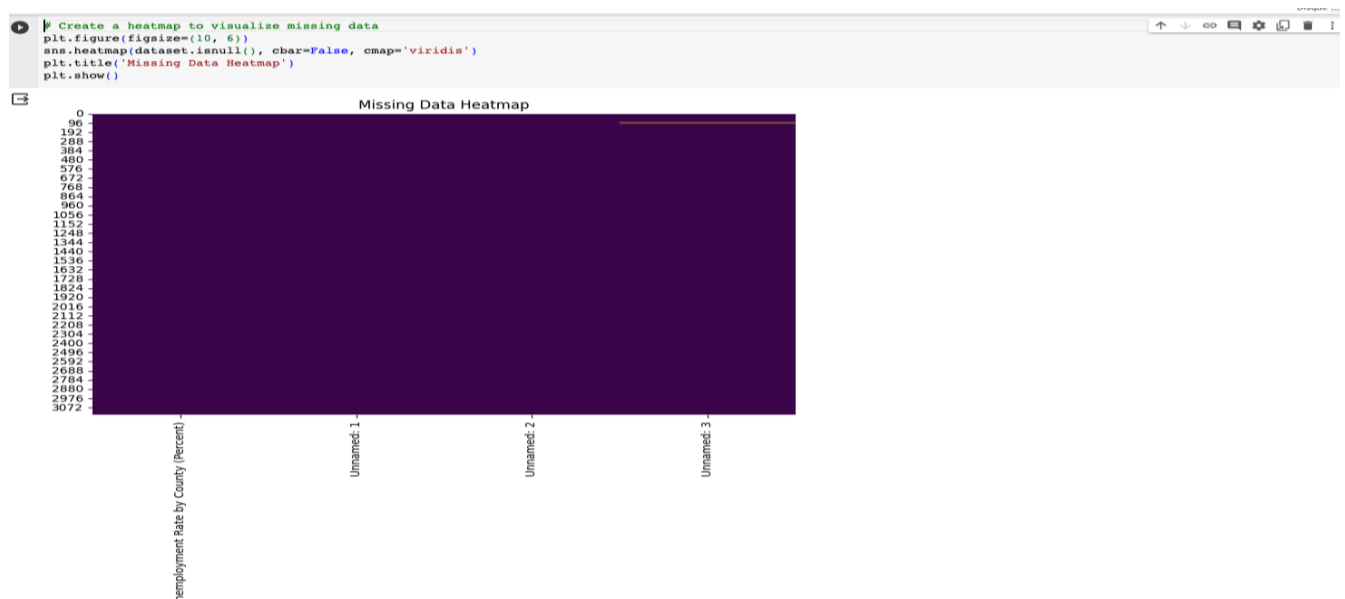
Positions of missing values in variable Unnamed: 3 : (array([ 92, 94, 1038, 2419, 2918]),)

[20] # Use the For loop to get missing values and get the positions in all the columns
for col_name in dataset.columns:
    print("position of the NaN values in all the columns", col_name, ":",np.where(pd.isna(dataset[col_name])))

position of the NaN values in all the columns 2023 June Unemployment Rate by County (Percent) : (array([], dtype=int64),)
position of the NaN values in all the columns Unnamed: 1 : (array([], dtype=int64),)
position of the NaN values in all the columns Unnamed: 2 : (array([], dtype=int64),)
position of the NaN values in all the columns Unnamed: 3 : (array([ 92, 94, 1038, 2419, 2918]),)

```

▾ Missing Map



Calculates and displays the percentage of missing values in each column using `round(dataset.isnull().sum() / len(dataset) * 100, 1)`.

```
[22] #Get the percentage of missing values in each column
missing_pct = round(dataset.isnull().sum()/len(dataset) * 100, 1)
print(missing_pct)

2023 June Unemployment Rate by County (Percent)    0.0
Unnamed: 1                                           0.0
Unnamed: 2                                           0.0
Unnamed: 3                                           0.2
dtype: float64
```

Data cleaning:

The code snippet below demonstrates the location-based replacement of values in the 'Unnamed: 3' column:

Cleaning the dataset

```
# Location based replacement
dataset.loc[92, 'Unnamed: 3'] = 7.51
dataset.loc[94, 'Unnamed: 3'] = 7.62
dataset.loc[1038, 'Unnamed: 3'] = 7.83
dataset.loc[2419, 'Unnamed: 3'] = 7.94
dataset.loc[2918, 'Unnamed: 3'] = 8.15

print(dataset)
print(dataset['Unnamed: 3'].unique())
```

	2023 June Unemployment Rate by County (Percent)	Series ID	Region Name
0		ALAUTA1URN	Autauga County, AL
1		ALBALD0URN	Baldwin County, AL
2		ALBARBSURN	Barbour County, AL
3		ALBIBB7URN	Bibb County, AL
4	
...	
3140		WYSWEE7URN	Sweetwater County, WY
3141		WYTETO9URN	Teton County, WY
3142		WYUINT1URN	Uinta County, WY
3143		WYWASH3URN	Washakie County, WY
3144		WYWEST5URN	Weston County, WY

	Unnamed: 2	Unnamed: 3
0	Region Code	01-06-2023
1	1001	2.3
2	1003	2.3
3	1005	5
4	1007	2.9
...
3140	56037	3.6
3141	56039	1.7
3142	56041	3.4
3143	56043	3.4
3144	56045	2.2

```
[3145 rows x 4 columns]
['01-06-2023' '2.3' '5' '2.9' '2.7' '3.2' '3' '2.6' '2.5' '4.3' '5.2'
 '2.4' '3.1' '3.9' '2.8' '2.2' '5.7' '3.7' '4.4' '3.8' '3.6' '2.1' '6.2'
 '3.5' '2' '4.5' '7.5' '3.3' '10.1' '7.2' '5.1' '4' '4.9' '8.1' '10' '6.4'
 '5.8' 7.51 '18.1' 7.62 '4.7' '6.3' '8.9' '8.8' '4.8' '5.5' '4.6' '7.3'
 '4.2' '14.3' '5.3' '4.1' '3.4' '7.1' '5.4' '12.2' '5.6' '6.5' '16.9'
 '8.4' '7.7' '9.5' '6' '5.9' '6.8' '7.6' '6.6' '1.7' '1.9' '1.8' '6.1'
 '8.5' '8.2' '9.4' '6.9' 7.83 '7.4' '7.8' '9.7' '9.1' '6.7' '10.2' '1.5'
 '1.6' '1.4' '7.9' '14.1' '9.8' '1.2' '9' '8' 7.94 '0.9' '0.4' '9.6' 8.15]
```

At row index 92, the value in the 'Unnamed: 3' column was replaced with 7.51.

At row index 94, the value in the 'Unnamed: 3' column was replaced with 7.62.

At row index 1038, the value in the 'Unnamed: 3' column was replaced with 7.83.

At row index 2419, the value in the 'Unnamed: 3' column was replaced with 7.94.

At row index 2918, the value in the 'Unnamed: 3' column was replaced with 8.15.

Following the replacement, the changed dataset was printed using `print(dataset)` to display the updated values. To check the changes, the unique values in the 'Unnamed: 3' column were displayed with `print(dataset['Unnamed: 3'].unique())`.

The provided code snippet demonstrates the process of replacing missing values:

The `.fillna(8.643, inplace=True)` method is applied to the 'Unnamed: 3' column of the 'dataset' DataFrame. This method replaces any missing values in the specified column with the numeric value 8.643. The `inplace=True` argument ensures that the changes are made directly to the 'dataset' DataFrame.

```

# Replace missing values with a number
file_path = '2023 June Unemployment Rate by County (Percent).csv'
dataset = pd.read_csv(file_path)
x = dataset.iloc[:, :-1].values
y = dataset.iloc[:, -1].values
dataset['Unnamed: 3'].fillna(8.643, inplace=True)
print(dataset)
print(dataset['Unnamed: 3'].unique())

```

Series ID	Region Name	Unnamed: 1 \
0	ALAUTA1URN	Autauga County, AL
1	ALBALD0URN	Baldwin County, AL
2	ALBARB5URN	Barbour County, AL
3	ALBIBB7URN	Bibb County, AL
4
3140	WYSWEE7URN	Sweetwater County, WY
3141	WYTETO9URN	Teton County, WY
3142	WYUINT1URN	Uinta County, WY
3143	WYWASH3URN	Washakie County, WY
3144	WYWEST5URN	Weston County, WY

Region Code	01-06-2023	Unnamed: 2	Unnamed: 3
0	1001	2.3	
1	1003	2.3	
2	1005	5	
3	1007	2.9	
4
3140	56037	3.6	
3141	56039	1.7	
3142	56041	3.4	
3143	56043	3.4	
3144	56045	2.2	

```

[3145 rows x 4 columns]
['01-06-2023' '2.3' '5' '2.9' '2.7' '3.2' '3' '2.6' '2.5' '4.3' '5.2'
 '2.4' '3.1' '3.9' '2.8' '2.2' '5.7' '3.7' '4.4' '3.8' '3.6' '2.1' '6.2'
 '3.5' '2' '4.5' '7.5' '3.3' '10.1' '7.2' '5.1' '4' '4.9' '8.1' '10' '6.4'
 '5.8' 8.643 '18.1' '4.7' '6.3' '8.9' '8.8' '4.8' '5.5' '4.6' '7.3' '4.2'
 '14.3' '5.3' '4.1' '3.4' '7.1' '5.4' '12.2' '5.6' '6.5' '16.9' '8.4'
 '7.7' '9.5' '6' '5.9' '6.8' '7.6' '6.6' '1.7' '1.9' '1.8' '6.1' '8.5'
 '8.2' '9.4' '6.9' '7.4' '7.8' '9.7' '9.1' '6.7' '10.2' '1.5' '1.6' '1.4'
 '7.9' '14.1' '9.8' '1.2' '9' '8' '0.9' '0.4' '9.6']

```

Data Export:

To store the cleaned and changed dataset for further analysis, the code uses `dataset.to_csv("new_dataset.csv", index=False)` to save the modified DataFrame to a new CSV file named 'new_dataset.csv'.

```

[53] # Write the DataFrame to a CSV file
dataset.to_csv("new_dataset.csv", index=False) # Set index=False to omit writing row numbers

```

Discuss possible problems you plan to investigate for future studies.

For future studies, I plan to investigate:

Exploratory Data Analysis (EDA): To acquire deeper insights, try undertaking more complete EDA, such as investigating connections between variables, displaying distributions, and detecting outliers.

Data Visualization: Extend the visualization capabilities to incorporate different sorts of plots, such as histograms, box plots, scatter plots, or time series plots, to expose new insights in the data.

Check whether Simpson's paradox exists in your dataset and explain what you find and why you choose these visualization methods.

Finally, this report highlights the critical procedures involved in preparing a dataset comprising unemployment rates by county. This function prepares the data for future analysis, visualization, or modeling by importing the data, investigating its features, addressing missing values, and exporting the cleaned dataset.